

Correspondence

A Survey of Decision Tree Classifier Methodology

S. Rasoul Safavian and David Landgrebe

Abstract—Decision tree classifiers (DTC's) are used successfully in many diverse areas such as radar signal classification, character recognition, remote sensing, medical diagnosis, expert systems, and speech recognition, to name only a few. Perhaps, the most important feature of DTC's is their capability to break down a complex decision-making process into a collection of simpler decisions, thus providing a solution that is often easier to interpret. A survey of current methods for DTC designs and the various existing issues are presented. After considering potential advantages of DTC's over single-state classifiers, the subjects of tree structure design, feature selection at each internal node, and decision and search strategies are discussed. Some remarks concerning the relation between decision trees and neural networks (NN) are also made.

I. INTRODUCTION

The decision tree classifier (DTC) is one of the possible approaches to multistage decision making; table look-up rules [34], decision table conversion to optimal decision trees [36], [43], [57], [94], and sequential approaches [22], [80] are others. The basic idea involved in any multistage approach is to break up a complex decision into a union of several simpler decisions, hoping the final solution obtained this way would resemble the intended desired solution. A complete review of multistage recognition schemes is given by Dattatreya and Kanal [14]. Some of the differences between the different multistage schemes are addressed in Kulkarni and Kanal [45]. The emphasis of this paper is on the hierarchical approaches. Hierarchical classifiers are a special type of multistage classifiers that allow rejection of class labels at intermediate stages.

The organization of the paper is as follows. Section II contains the preliminaries, definitions, and terminologies needed for later sections; Section III explains the motivations behind DTC's and their potential use and drawbacks; Section IV addresses the problems of tree structure design, feature selection, and decision rules to be used at each internal node. The tree structure design is considered from top-down, bottom-up, hybrid, and tree growing-pruning points of view. In considering the feature subset selection at the internal nodes of the tree, a newly proposed approach based on application of neural networks is considered. Some possible search strategies are also mentioned in this section. These problems are addressed from the Bayesian, the decision-theoretic, and the information-theoretic and entropy reduction approaches. Other issues such as incremental tree design, the tree capability to generalize, the missing data value problem, the robustness of tree design, and the relation between decision trees and neural networks are discussed in Section V. Final comments and conclusions are provided in Section VI.

Manuscript received October 21, 1990; revised December 22, 1990. This work was supported in part by the National Science Foundation under Grant NSF ECS8507405 and by the National Aeronautics and Space Agency under Grant NAGW-925.

The authors are with the School of Electrical Engineering, Purdue University, West Lafayette, IN 47907.
IEEE Log Number 9042702.

II. PRELIMINARIES

We briefly describe some necessary terminology for describing trees (see Aho *et al.* [3] for more details).

Definitions

- 1) A graph $G = (V, E)$ consists of a finite, nonempty set of nodes (or vertices) V and a set of edges E . If the edges are ordered pairs (ν, w) of vertices, then the graph is said to be *directed*.
- 2) A path in a graph is a sequence of edges of the form $(\nu_1, \nu_2), (\nu_2, \nu_3), \dots, (\nu_{n-1}, \nu_n)$. We say the path is from ν_1 to ν_n and is of the length n .
- 3) A directed graph with no cycles is called a *directed acyclic graph*. A *directed* (or *rooted*) *tree* is a directed acyclic graph satisfying the following properties:
 - a) There is exactly one node, called the *root*, which no edges enter. The root node contains all the class labels.
 - b) Every node except the root has exactly one entering edge.
 - c) There is a unique path from the root to each node.
- 4) If (ν, w) is an edge in a tree, then ν is called the *father* of w , and w is a *son* of ν . If there is a path from ν to w ($\nu \neq w$), then ν is a proper *ancestor* of w and w is a proper *descendant* of ν .
- 5) A node with no proper descendant is called a *leaf* (or a *terminal*). All other nodes (except the root) are called *internal* nodes.
- 6) The *depth* of a node ν in a tree is the length of the path from the root to ν . The *height* of node ν in a tree is the length of a largest path from ν to a leaf. The *height* of a tree is the height of its root. The *level* of a node ν in a tree is the height of the tree minus the depth of ν .
- 7) An *ordered tree* is a tree in which the sons of each node are ordered (normally from left to right).
- 8) A *binary tree* is an ordered tree such that:
 - a) each son of a node is distinguished either as a *left son* or as a *right son*;
 - b) no node has more than one left son nor more than one right son.
- 9) The balance of a node ν in a binary tree is $(1 + L)/(2 + L + R)$, where L and R are the number of nodes in the left and right subtrees of ν . A binary tree is α -balanced, $0 < \alpha \leq 1$, if every node has balance between α and $1 - \alpha$. A $1/2$ -balanced binary tree is also known as a *complete* tree.

We will denote any tree by T . Fig. 1 shows an example of a general decision tree.

- 10) If two internal nodes contain at least one common class, then it is said that the nodes have *overlap* classes. See Fig. 2.
- 11) The average number of layers from the root to the terminal nodes is referred to as the *average depth* of the tree. The average number of internal nodes in each level of the tree is referred to as the *average breadth* of the tree. In general, the average breadth of the tree will reflect the relative weight given to classifier accuracy whereas the average depth of the tree will reflect the weight given to efficiency [77].

Let us also introduce the following. Let (\underline{X}, Y) be jointly distributed random variables with q -dimensional vector \underline{X} denoting a





Fig. 1. Example of a general decision tree.

Fig. 2. Part of the tree with overlap at nodes i and k ; i.e., classes 3 and 5 are repeated.

pattern or feature vector and \mathbf{Y} denoting the associated class label of \mathbf{X} . The components of \mathbf{X} are the features.

Definitions

- 12) \mathbf{X} is called an *ordered or numerical* (see Breiman *et al.* [7]) pattern if its features take values from an ordered set, and *categorical* if its features take values from a set not having a natural ordering. Ordered or numerical features can have either discrete or continuous values. For simplicity of notation let us assume that \mathbf{X} is of continuous ordered type; furthermore, let \mathbf{X} take values from R^q . Let \mathbf{Y} take integer values $\{1, 2, \dots, J\}$; i.e., there are J classes of concern. Then the goal of any classification scheme in general, and DTC in particular, is to estimate \mathbf{Y} based on observing \mathbf{X} .
- 13) A decision rule $d(\cdot)$ is a function that maps R^q into $\{1, 2, \dots, J\}$ with $d(\mathbf{X})$ representing the class label of feature vector \mathbf{X} . The *true misclassification rate* of d , denoted by $R^*(d)$, is

$$R^*(d) = p(d(\mathbf{X}) \neq \mathbf{Y}) \quad (1)$$

where $p(\cdot)$ denotes the probability.

Let us denote the available labeled samples as $\mathcal{L} = \{(\mathbf{X}_n, \mathbf{Y}_n), n = 1, 2, \dots, N\}$.

- 14) Usually the available labeled samples are divided into two parts: the training sample $\mathcal{L}^{(1)}$ and the test sample $\mathcal{L}^{(2)}$. Commonly, $\mathcal{L}^{(1)}$ is taken to be randomly sampled 2/3 of \mathcal{L} and $\mathcal{L}^{(2)}$ the remaining 1/3 of \mathcal{L} .
- 15) Due to the difficulty of computing the true misclassification rate $R^*(d)$, it is usually estimated from either the training set or the test set. Let us denote the estimated misclassification rate as $R(d)$. When the training set is used to estimate $R^*(d)$, $R(d)$ is called the *resubstitution estimate* of $R^*(d)$, and when

the test sample is used to estimate $R^*(d)$, $R(d)$ is called a *test sample estimate* of $R^*(d)$. In either case, the misclassification rate is simply estimated by the ratio of samples misclassified to the total number of samples used to do the test. A more complex method of estimating the misclassification rate is the K -fold cross-validation method. Here, the data set \mathcal{L} is divided into k nearly equal parts $\mathcal{L}^1, \mathcal{L}^2, \dots, \mathcal{L}^K$ (usually $K = 10$). Then $\mathcal{L} - \mathcal{L}^k$ is used for the training sample, and \mathcal{L}^k is used for the test sample. Next the test sample estimate of the misclassification rate for each k is found, and averaged over K . Note that when $K = N$ (the size of the labeled sample), N -fold cross-validation is also called the "leave-one-out" estimate.

III. POTENTIALS AND PROBLEMS WITH DCT's

DTC's are attractive for the following reasons [40], [52], [53], [59], [77], [81]:

- 1) Global complex decision regions (especially in high-dimensional spaces) can be approximated by the union of simpler local decision regions at various levels of the tree.
- 2) In contrast to conventional single-stage classifiers where each data sample is tested against all classes, thereby reducing efficiency, in a tree classifier a sample is tested against only certain subsets of classes, thus eliminating unnecessary computations.
- 3) In single-stage classifiers, only one subset of features is used for discriminating among all classes. This feature subset is usually selected by a globally optimal criterion, such as maximum average interclass separability [77]. In DTC's, on the other hand, one has the flexibility of choosing different subsets of features at different internal nodes of the tree such that the feature subset chosen optimally discriminates among the classes in that node. This flexibility may actually provide performance improvement over a single-stage classifier [58], [97].
- 4) In multivariate analysis, with large numbers of features and classes, one usually needs to estimate either high-dimensional distributions (possibly multimodal) or certain parameters of class distributions, such as *a priori* probabilities, from a given small sized training data set. In so doing, one usually faces the problem of "high-dimensionality." This problem may be avoided in a DTC by using a smaller number of features at each internal node without excessive degradation in the performance.

The possible drawbacks of DTC, on the other hand, are as follows.

- 1) Overlap (for the definition, see Section II), especially when the number of classes is large, can cause the number of terminals to be much larger than the number of actual classes and thus increase the search time and memory space requirements. Possible solutions to this problem will be addressed in Section IV.
- 2) Errors may accumulate from level to level in a large tree. It is pointed out by Wu *et al.* [95] that one cannot simultaneously optimize both the accuracy and the efficiency; for any given accuracy a bound on efficiency must be satisfied.
- 3) Finally, there may be difficulties involved in designing an optimal DTC. The performance of a DTC strongly depends on how well the tree is designed.

IV. DESIGN OF A DTC

The main objectives of DTC's are: 1) to classify correctly as much of the training sample as possible; 2) generalize beyond the training sample so that unseen samples could be classified with as high of an accuracy as possible; 3) be easy to update as more training sample becomes available (i.e., be incremental—see Section IV); 4) and have

as simple a structure as possible. Then the design of a DTC can be decomposed into following tasks [45], [49], [50]:

- 1) the appropriate choice of the tree structure,
- 2) the choice of feature subsets to be used at each internal node,
- 3) the choice of the decision rule or strategy to be used at each internal node.

When a Bayes point of view is pursued, the optimal tree design may be posed as the following optimization problem:

$$\min_{T, F, d} p_e(T, F, d) \quad (2)$$

subject to: Limited training sample size

where P_e is the overall probability of error, T is a specific choice of the tree structure, F and d are the feature subsets and decision rules to be used at the internal nodes, respectively. The implication of the aforementioned constraint is that, with a limited training sample size, the accuracy of the estimates of class conditional densities may deteriorate as the number of features increases. This is also known as the Hughes phenomena [38]. In practice this places a limit on the number of features that may be utilized. That is, out of say L , available features, one forms a new feature set of size N , e.g., by subset selection or by combining features, where usually $N \ll L$. Note that for a fixed size N , selected as the best feature subset to discriminate among *all* the classes, the minimum probability of error decision rule, by definition, is the single-stage Bayes rule. But, if one is allowed to select different feature subsets for differentiating between different groups of classes (i.e., a tree structure), one may be able to obtain even smaller probabilities of error than those predicted by Bayes decision rule.

The previous optimization problem can be solved in two steps [50]:
Step 1: For a given T and F , find $d^* = d^*(T, F)$ such that

$$p_e(T, F, d^*(T, F)) = \min_d p_e(T, F, d). \quad (3)$$

Step 2: Fine T^* and F^* such that

$$p_e(T^*, F^*, d^*(T^*, F^*)) = \min_{T, F} p_e(T, F, d^*(T, F)). \quad (4)$$

It should be noted here that no mention of time complexity or computation speed has been made so far. Including these factors would make the optimization problem even more difficult. Swain and Hauska [77], and Wang and Suen [81]–[84] have offered some ways to include time efficiency (i.e., speed) into the analysis. These are discussed in the latter part of the paper.

When the information theoretic point of view is pursued, however, the optimal design of the tree involves maximizing the amount of average mutual information gain at each level of the tree. We will return to this point shortly.

A. Decision Tree Structure Design

Several methods [4], [6], [8], [15], [17], [20], [32], [37], [42], [45], [48], [53], [61], [64], [65], [68], [73], [77], [82], [95], [97], [98] have been proposed for the tree structure design. Some have no claim of optimality and utilize the available *a priori* knowledge for the design [4], [30], [52], [68], [74], [81], [95], [97], while others apply mathematical programming methods such as dynamic programming or branch-and-bound techniques [45], [46], [48], [57], [61]. Some studies [53], [65], [71], [73] have attempted to combine the tree structure design task with feature selection and the decision rules to be used throughout the tree.

Some of the common optimality criteria for tree design are: minimum error rate, min-max path length, minimum number of nodes in the tree, minimum expected path length, and maximum average mutual information gain. The optimal tree is constructed




Fig. 3. Two-dimensional feature space showing two possible partitionings of the same space, each providing 100% classification accuracy. Samples from the three classes are designated x , o , and $+$. The two sets of decision boundaries are designated by the solid lines and the dashed lines.




Fig. 4. Decision trees for the feature space partitions of Fig. 3. (a) Decision tree for the solid line decision boundaries of Fig. 3. (b) Decision tree for the dashed line decision boundaries of Fig. 3.

recursively through the application of various mathematical programming techniques such as dynamic programming [56] with look-ahead (or back) capabilities to approach global optimality. A basic problem with any of the suggested optimization methods is their computational infeasibility; usually large amounts of computational time and storage (space) are required. This in practice has placed a restriction on the type of tree and/or the number of features (variables) to be used at each node [97].

In many pattern recognition problems at the outset, one has only a set of labeled samples \mathcal{L} from the set of possible classes. The problem is then, given \mathcal{L} , find a decision tree T that optimizes some cost function, such as the average number of nodes in the tree. However, there may be several ways the labeled sample space could be partitioned, all being equally attractive and acceptable. For instance, Fig. 3 shows two possible binary partitionings using hyperplanes perpendicular to the feature axes of the space with three classes and two features; both provide 100% correct classification of the labeled samples.

For each partitioning, there is a corresponding binary tree. Fig. 4 shows the binary trees corresponding to the partitioning of Fig. 3. The goal, is to construct a binary tree that is equivalent in the sense of giving the same decision, but optimal in the sense of having the minimum average number of nodes [44], or more generally satisfying some size optimality criterion. Meisel and Michalopoulos [57], and Payne and Meisel [61] address exactly this problem.

The algorithms they provide are basically composed of two parts. In the first part, they find a sufficient partitioning of the space; there are actually many ways to do this [2], [24], [61], [73]. A general




Fig. 5. A regular lattice constructed by extending the original partitioning, with each element of the lattice assigned a class label.

approach [61] is to sequentially examine the univariate projections of the sample onto each feature axes, then find the hyperplanes perpendicular to the feature axes that partitions the projected samples as accurately as possible, with 100% accuracy if samples are linearly separable. These hyperplanes partition the space into hypercuboids. Also, the problem is then transformed to one with discrete features. Each hypercuboid is labeled with the class label of the majority of the samples in that subregion. The hyperplanes are extended to infinity to create a regular lattice, L , thus each lattice point has a class label and a feature subset associated with it. See Fig. 5. Lattice L can be binary partitioned on the feature axis i into mutually exclusive left and right sets, where a left (right) set includes lattice elements with feature i values smaller (larger) than the threshold value. A binary partition can also be denoted by a five-tuple (n, i, m, n^L, n^R) , where n denotes a decision node label for the partition, i and m are the feature axis and the threshold value used for the partition, respectively, and n^L and n^R are the node labels for partition of the left and right sets, respectively. Note that a binary tree is a collection of nested binary partitions, and can conveniently be represented in the following recursive form

$$T = \left\{ (n, i, m, n^L, n^R), T^L, T^R \right\} \quad (5)$$

starting from the root; where T^L and T^R denote the subtrees defined on the left and right sets of a partition. For the binary trees defined on this lattice L , the general cost function considered has the form [61]:

$$C(T) = F(L, C(T^L), C(T^R)) \quad (6)$$

where $F(L, \cdot, \cdot)$ is componentwise nondecreasing and $F \geq 0$. The optimization problem, is then given L , find T^* such that

$$C^*(L) = \min_T C(T) = C(T^*). \quad (7)$$

The optimal tree, therefore, could be recursively constructed through the application of invariant embedding (dynamic programming).

Remarks

- 1) Only one feature is examined at each node.
- 2) The algorithms are feasible only for a small number of features, else the size of the lattice becomes large and storage space requirements become a problem. This is because during the optimization process intermediate results must be fully accessible. This perhaps is the main limitation.

The problem of designing a truly optimal DTC seems to be a very difficult problem. In fact it has been shown by Hyafil and Rivest [39] that the problem of constructing optimal binary trees, optimal in the

sense of minimizing the expected number of tests required to classify an unknown sample is an NP -complete problem and thus very unlikely of non-polynomial time complexity. It is conjectured that more general problems, i.e., the problem with a general cost function or minimizing the maximum number of tests (instead of average) to classify an unknown sample would also be NP -complete. They also conjecture that no sufficient algorithm exists (on the supposition that $P \neq NP$) and thereby supply motivation for finding efficient heuristics for constructing near-optimal decision trees.

The various heuristic methods for construction of DTC can roughly be divided into four categories: *Bottom-Up* approaches, *Top-Down* approaches, the *Hybrid* approach and tree *Growing-Pruning* approaches. In a bottom-up approach [52], a binary tree is constructed using the training set. Using some distance measure, such as Mahalanobis-distance, pairwise distances between *a priori* defined classes are computed and in each step the two classes with the smaller distance are merged to form a new group. The mean vector and covariance matrix for each group are computed from the training samples of classes in that group, and the process is repeated until one is left with one group at the root. This has some of the characteristics of an unsupervised hierarchical clustering approach. In a tree constructed this way, the more obvious discriminations are done first, near the root, and more subtle ones at later stages of the tree. It is also recommended [52] that from a processing speed point of view, the tree should be constructed such that the most frequently occurring classes are recognized first.

It is worth noting that usually most decision tree designs are restricted to having a binary structure. This is not really a restriction since any *ordered* tree can be uniquely transformed into an *equivalent* binary tree [65].

In top-down approaches, the design of a DTC reduces to the following three tasks:

- 1) the selection of a node splitting rule,
- 2) the decision as to which nodes are terminal,
- 3) the assignment of each terminal node to a class label.

Of the previous three tasks, the class assignment problem is by far the easiest. Basically, to minimize the misclassification rate, terminal nodes are assigned to the classes that have the highest probabilities. These probabilities are usually estimated by the ratio of samples from each class at that specific terminal node to the total number of samples at that specific terminal node. This is the basic majority rule; i.e., assign to the terminal node the label of the class that has most samples at that terminal node.

It seems that most of the research in DTC design has concentrated in the area of finding various *splitting rules*; this also naturally embodies the *termination rules*. A summary of some of the research done in this direction follows. Wu *et al.* [95] have suggested that a histogram based on an interactive top-down approach for the tree design may be especially useful for remote sensing applications. In this approach, the histogram of the training data of all classes is plotted on each feature axis with the same scale. By observing the histograms, a threshold is selected to partition those classes into several groups. If a group contains more than one class, the same procedure is repeated until each group contains only one class. The drawback of this method is that only a few features (usually one) are used at each stage. Interaction between features cannot be observed.

When the distributions of classes under consideration are known, You and Fu [97] proposed the following approach for design of a *linear* binary tree classifier. Since even for moderately small numbers of features and classes, the number of possible trees is astronomically large, they suggested two restrictions to reduce the size of the search space. First, limit the number of features to be used at each stage. Secondly, for the sake of accuracy, specify

tolerable error probabilities at each stage. Obviously the choice of linear classifiers and a binary tree structure is made to decrease computational complexity and time and thus to increase the speed. Again, using a distance measure, such as Bhattacharyya distance, classes at each node are divided into two groups. Then, using an iterative procedure with an initial guess, a classifier is found that provides minimum probability of error. If this error exceeds the pre-assigned error bound, the class that commits the maximum error is taken from consideration (i.e., included in both of the following subgroups causing overlap) and, repeating the previous procedure, a new classifier is found and the error assessed. Presuming that now the error is below the specified error bound, new subgroups are formed, else the entire aforementioned process is repeated. Assuming that the error calculation function is first-order differentiable with respect to the coefficients of the linear equations of the classifier, they recommend the Fletcher–Powell algorithm [19].

Remark: Even though allowing overlap is one way of improving recognition rate, caution must be exercised since overuse of it will cause the number of terminals to be much larger than the actual number of classes, thus reducing the efficiency of the tree classifier. This could be a serious drawback in *large-class* problems.

A heuristic procedure that incorporates the time component in the tree structure design is proposed by Swain and Hauska [77], and Wu *et al.* [95]. For every node n_i an evaluation function $E(n_i)$ is defined as

$$E(n_i) = -T(n_i) - W \cdot e(n_i) + \sum_{j=1}^{C_i} p_{i+j} \cdot E(n_{i+j}) \quad (8)$$

where $T(n_i)$ and $e(n_i)$ are the computation time and the classification error associated with node n_i , respectively, W is the weighting factor, specified by the user, reflecting the relative importance of accuracy to the computation time, and C_i is the number of descendent nodes of n_i . The third term on the right-hand side of (8) is the sum of the evaluation functions at the descendent nodes of n_i , weighted by their probabilities of access from n_i . Obviously in this forward top–down search procedure, the configuration of the descendent nodes of n_i is not known yet. Their evaluation function values, however, can be lower bound estimated by the values of the corresponding evaluation functions if the usual single-stage (one shot) classifier were used at n_i . The configuration that provides the maximum evaluation function among all the candidate configurations is selected and the process is repeated at the following nodes. As pointed out in [77], this node-by-node design approach, using only local (not global) information, can provide at best only a suboptimal result.

Since in top–down approaches to tree design, sets of classes are successively decomposed into smaller subsets of classes, some researchers have attempted to combine the tree design problem and feature selection problem by using an appropriate splitting criteria at each internal node [53], [65], [71], [73]. Rounds [65] has suggested Kolmogorov–Smirnov (K–S) distance and test as the splitting criteria. Considering a two-class problem with a binary tree structure in mind, only one feature is selected at each internal node, with the corresponding threshold value. The rationale of this approach is that crossover points of probability distributions of two classes (especially when classes have multimodal distributions) are good locations to place the partitioning planes. These crossover points also correspond to the relative maxima of difference of cumulative distributions, also known as the Kolmogorov–Smirnov (K–S) distance. In the absence of true class distributions, empirical distributions are estimated from the training set. The K–S distance being a random variable, a measure of confidence is provided by the K–S test. At each node for each feature, the K–S distances and confidence level are estimated. The actual value of the feature that

provides more than two maxima is used as the threshold value to create the descendant nodes.

This approach provides a means for evaluating the discrimination power and significance of each feature in the classification process but has the following shortcomings. Firstly, it requires as many decision trees to be developed as there are pattern classes [73]. Second, this method uses only one feature at each node. Third, even for one feature and for a two-class problem, it results in a tree with differentiated structure and a large number of nodes. Furthermore, the K–S distance depends on the sample sizes of both classes at each internal node, which would be in general different at the different internal nodes. The distribution of this statistic under certain restrictions on the relative values of sample sizes for each class can be found in Gnedenko and Kareljuk [26] and Rounds [66].

Li and Dubes [53] proposed a permutation statistic as the splitting criterion. The permutation statistic measures the degree of similarity between two binary vectors: the vector of thresholded feature values, obtained from the training data, and the vector of known pattern class labels. The advantages of this splitting criterion over, for instance, a statistic based on distance between empirical distributions (e.g., K–S distance) are its known distribution and its independence from the number of training patterns. Even though the permutation statistic is based on the two-class assumption, multiclass problems can be handled by a sequence of binary decision trees.

In general, the basic idea in choosing *any* splitting criteria at an internal nodes is to make the data in the son nodes “purer.” One way to accomplish this task [7] is to define an *impurity function* $i(t)$ at every internal node t . Then suppose that for the internal node t there is a candidate split S that divides node t into left son t_L and right son t_R such that a proportion p_L of the cases in t go into t_L and a proportion p_R go into t_R . One could then define the goodness of the split S to be the decrease in impurity:

$$\Delta i(S, t) = i(t) - i(t_L)p_L - i(t_R)p_R. \quad (9)$$

Then choose a split that minimizes $\Delta i(S, T)$ over all splits S in some set S . The impurity function suggested in [7] is the *Gini* index defined as

$$i(t) = \sum_{i \neq j} p(i/t)p(j/t) \quad (10)$$

where $p(i/t)$ is just the probability of a random sample \underline{X} belonging to the class i , given we are at node t .

The final component in top–down DTC design is to determine when the splitting should be stopped; i.e., we need a stopping rule. Early approaches to selecting terminal nodes were of the form: set a threshold $\beta > 0$ and declare a node t terminal if

$$\max_{S \in \mathcal{S}} \Delta i(S(t), t) < \beta. \quad (11)$$

The problem with this rule is that usually partitioning is halted too soon at some nodes and too late at some others. The other major problem in deciding when to terminate splitting is the following. Suppose we have classes whose underlying distributions overlap; that is the true misclassification rate is positive ($R^*(d) > 0$). Furthermore, suppose we have a training set \mathcal{L} , and we use a splitting approach (e.g., impurity function) along with a terminal labeling approach mentioned earlier, and continue the splitting until every node has pure samples; i.e., samples from only one class. Now consider the resubstitution estimate (see Section II, definition 15) of the misclassification rate of the tree T defined as

$$R(T) = \sum_{t \in \hat{T}} r(t)p(t) = \sum_{t \in \hat{T}} R(t) \quad (12)$$

where $r(t)$ is the resubstitution estimate of the misclassification rate, given a sample falls into node t , $p(t)$ is the probability of a random sample falling into node t , and \hat{T} is the set of terminal nodes of the tree T . Notice that for previous scenario $R(T) = 0!$. Obviously this estimate is very biased. That is, in general $R(T)$ decreases as the number of terminal nodes increases. But the tree so constructed usually fails miserably in classification of the test data. Breiman *et al.* [7] conjectured that decision tree design is rather insensitive to a variety of splitting rules, and it is the stopping rule that is crucial. Breiman *et al.* [7] suggest that instead of using stopping rules, continue the splitting until all the terminal nodes are pure or nearly pure, thus resulting in a large tree. Then selectively prune this large tree, getting a decreasing sequence of subtrees. Finally, use cross-validation to pick out the subtree that has the lowest estimated misclassification rate. To go any further, we need the following definitions (see [7] for more details).

Definitions

- 16) A branch T_t of tree T with root node $t \in T$ consists of node t and all the descendants of t in T .
- 17) Pruning a branch T_t from a tree T consists of deleting from T all descendants of t . Denote the pruned tree as $T - T_t$.
- 18) If by successively pruning off branches from T we obtain the subtree T' , we call T' a pruned subtree of T and denote it by $T' < T$.

Recall that the size of a tree is of utmost importance. In order to include the complexity of tree T in the design process, define a cost-complexity measure $R_\alpha(T)$ as

$$R_\alpha(T) = R(T) + \alpha |\hat{T}| \quad (13)$$

where $R(T)$ has the usual meaning of estimated misclassification rate of tree T , and $\alpha \geq 0$ is the complexity cost per terminal node. That is, the complexity of tree T is measured by the number of its terminal nodes. Let us denote the fully grown tree as T_{\max} . Then the objective of a pruning process, for every α , is to find the subtree $T(\alpha) \leq T_{\max}$ such that

$$R_\alpha(T(\alpha)) = \min_{T \leq T_{\max}} R_\alpha(T). \quad (14)$$

Notice that as α is varied from sufficiently small values where T_{\max} itself is the optimally pruned tree, to sufficiently large values where T_{\max} is pruned so far that only the root node $\{\text{root}\}$ is left, one obtains a finite sequence of trees $T_1, T_2, \dots, \{\text{root}\}$; where T_i corresponds to the optimal subtree for $\alpha = \alpha_i$. Then among these trees, the tree that provides the smallest test sample or cross-validation estimate of the misclassification rate is selected. (See [7] for more details.)

Note that the pruning algorithm is essentially a (multipass) top-down algorithm. The tree growing and pruning of Breiman *et al.* [7] has been incorporated into a computer program known as classification and regression trees (CART). Following are some of the shortcomings of CART. First of all, CART allows only either a single feature or a linear combination of features at each internal node. Second, CART is computationally very expensive as it requires generation of multiple auxiliary trees. Finally and perhaps most importantly, CART selects the final pruned subtree from a parametric family of pruned subtrees, and this parametric family may not even include the optimal pruned subtree. There are, of course, other tree growing and, more importantly, pruning algorithms (e.g., [12], [25]). Gelfand *et al.* [25] propose an *iterative* tree growing and pruning algorithm based on the following ideas.

- 1) Divide the data set into two approximately equal sized subsets and iteratively grow the tree with one subset and prune it with other subset.
- 2) Successively interchange the role of the two subsets.

Their pruning algorithm is a simple and intuitive *one-pass* bottom-up approach that starts from the terminal nodes and proceeds up the tree, pruning away branches. They prove the convergence of their algorithm, and their experimental results on waveform recognition, along with the theoretical proof, suggests the superiority of their algorithm to the pruning algorithm of Breiman *et al.* [7]. (See [25] for details.)

Recently, a *hybrid* method was proposed by Kim and Landgrebe [41] that uses both bottom-up and top-down approaches sequentially. The rationale for the proposed method is that in a top-down approach such as hierarchical clustering of classes, the initial cluster centers and cluster shape information are unknown. It is also well known that the proper choice of initial conditions could considerably influence the performance of a clustering algorithm (e.g., speed of convergence and final clusters). These information can be provided by a bottom-up approach. The procedure is as follows. First, considering the entire set of classes, one uses a bottom-up approach such as [52] to come up with two clusters of classes. Then one computes the mean and covariance for each cluster and uses these information in a top-down algorithm such as [95] to come up with two new clusters. Every cluster is checked to see if it contains only one class. If so, that cluster is labeled as terminal; else the previous procedure is repeated. The procedure terminates when all the clusters are labeled as terminals. See [41] for details.)

In all the foregoing discussions, at each node a decision is made and only one path is traversed from the root of the tree to a terminal node. This is referred to as a "hard-decision system" by Schuerman and Doster [68]. In contrast to hard-decision systems, they propose an approach where all the *a posteriori* probabilities are approximated at each internal node without making a decision at any of those points; a decision is made only at the terminal nodes of the tree by selecting the maximum *a posteriori* probability. In other words, in a manner similar to single-stage approaches [23], *a posteriori* probabilities for the different classes are estimated, but in a sequence of steps. Obviously, one would expect to obtain the same *a posteriori* probabilities at the terminal nodes as the ones the global single stage classifier would give, assuming the estimated *a posteriori* probabilities at each internal node are the same as the actual values.

In a single-stage classifications, one way to classify an unknown sample into one of m classes is to evaluate $(m - 1)$ discriminant functions [18], [23]. When Bayes minimum probability of error is the desired criterion, these discriminant functions are the *a posteriori* probabilities. In a sequential approach to evaluate these probabilities, the conditional probabilities of classes at each of the output branches of an internal node are estimated, based on the set of classes and the full feature measurement vector at the input to that node. Obviously the structure of the tree and the accuracy of these estimates are somewhat related. But, clearly, if the true *a posteriori* probabilities could be furnished, the explicit structure of the tree would not be critical. So, at least conceptually, the input classes at each node can be decomposed in a way that *a posteriori* probabilities of the classes at the outgoing branches may be estimated with "maximal confidence." Again due to the enormous number of possible combinations at each node, the actual tree is usually constructed using training data and heuristic methods based on a separability measure.

Mean square polynomial discriminant functions, using training sets, are suggested [96] for estimating the *a posteriori* probabilities of classes at the output of each node branch given the set of classes at the input of the node and the full feature measurement vector.

Remarks

- 1) By thresholding these *a posteriori* probabilities at each node, if desired, the soft-decision approach can be converted to the conventional hard-decision approach.

- 2) When the number of classes is large and no overlap is allowed in the tree, the hard-decision approach raises the possibility of error accumulation [81]–[83]. Even though using heuristics such as proper tree search strategies (see Section IV-E) this error can be reasonably controlled [81], but the soft-decision approach avoids the problem altogether.
- 3) When the number of classes and features are large and the class distributions have significant overlap, a hard-decision approach usually covers up ambiguities by a “forced recognition” [68]. In the soft-decision approach, no decision is made until the last stage of the tree.
- 4) In general, however, the computational time complexities of the soft-decision approach may be a limiting factor.

B. Entropy Reduction and Information-Theoretic Approaches

A different point of view about pattern recognition is taken by Watanabe [85]–[93] and Smith and Grandy [75]. Since organization and structure of objects or a set of stochastic variables could be expressed in terms of entropy functions [86], Watanabe refers to the problem of pattern recognition as that of “seeing a form” or structure in an object [86], and he suggests ways to cast the problems of learning [88], [89], such as feature (variable) selection, dimensionality reduction [91], and clustering [90], [92], to mention a few, in terms of minimizing properly defined entropy functions. As noted by Suen and Want [76], [81] this point of view could be very attractive when the number of classes is large and calculation of Bayes error probabilities is not so simple. Basically, since at the root of a tree, a given sample could belong to any of the classes, the uncertainty is maximum. At the terminal nodes of the tree, the sample is eventually classified and the uncertainty is eliminated. So, an objective function for a tree design could be to minimize uncertainty from each level to the next level, or in other words, maximize entropy reduction at each stage. Since it is also desirable to have as few overlaps as possible, Suen and Wang [76] suggested an interactive iterative clustering algorithm (ISOETRP) with an objective function directly proportional to the entropy reduction and inversely proportional to some function of class overlap at each stage of the tree. Even though any entropy measure could be used, Shannon’s entropy, defined as

$$H = - \sum_i p_i \log p_i \quad (15)$$

where p_i is the *a priori* probability of class i , is preferred because of its strong additivity property [1].

Sethi and Sarvarayudu [73] propose a simple, yet elegant, method for the hierarchical partitioning of the feature space. The method is non-parametric and based on the concept of average mutual information. More specifically, let the average mutual information obtained about a set of classes C_k from the observation of an event X_k , at a node k in a tree T be defined as

$$I_k(C_k; X_k) = \sum_{C_k} \sum_{X_k} p(C_{ki}, X_{kj}) \cdot \log_2 \left[\frac{p(C_{ki}/X_{kj})}{p(C_{ki})} \right]. \quad (16)$$

Event X_k represents the measurement value of a feature selected at node k and has two possible outcomes; measurement values greater or smaller than a threshold associated with that feature at that node.

Then, the average mutual information between the entire set of classes, C , and the partitioning tree, T , can be expressed as

$$I(C; T) = \sum_{k=1}^L p_k \cdot I_k(C_k; X_k) \quad (17)$$

where p_k is the probability of the class set C_k and L is the number of internal nodes in the tree T .

The probability of misclassification, p_e , of a DTC T and the average mutual information $I(C; T)$ are also related as [73]

$$I(C; T) \geq - \sum_{j=1}^m [p(C_j) \cdot \log_2 p(C_j)] + p_e \cdot \log_2 p_e + (1 - p_e) \cdot \log_2 (1 - p_e) + p_e \cdot \log_2 (m - 1) \quad (18)$$

with equality corresponding to the minimum required average mutual information for a prespecified probability of error. Then a goal for design of the tree could be to *maximize* the average mutual information gain (AMIG) at *each* node k . The algorithm terminates, when the tree average mutual information, $I(C; T)$, exceeds the required minimum tree average mutual information specified by the desired probability of error. (See [73] for more details.) An alternative stopping criterion proposed by Talmon [78] is to test the statistical significance of the mutual information gain that results from further splitting a node.

A tree constructed by this top–down algorithm would have a binary structure, and since at each node the partitioning hyperplane (i.e., the optimum feature and its corresponding threshold) is selected to maximize the average mutual information at that node, the performance is optimized in a “local sense.” Once the partitioning is specified, the globally optimum tree could be obtained using the algorithm suggested by Payne and Meisel [61].

Chou and Gray [11] view decision trees as variable-length encoder–decoder pairs, and compare the performance of the decision trees to the theoretical limits given by the rate-distortion function. More specifically, they show that rate is equivalent to the average tree depth while distortion is the probability of misclassification. They also use greedy algorithms to design trees and compare their performances to certain classes of optimal trees under certain restrictive conditions. The applicability of their method is limited, in practice, to only small problems. Goodman and Smyth [29] prove that a tree designed based on a top–down mutual information algorithm is equivalent to a form of Shannon–Fano prefix coding. They drive several fundamental bounds for mutual information-based recursive tree design procedures. They also suggest a new stopping rule that they claim to be more robust in the presence of noise.

When the measurement variables (features) are discrete, as for instance, in computer diagnostics of diseases or in stroke directions or number of line crossings in character recognition, the conventional classifier design or feature selection techniques face practical difficulties [71]. This is mainly due to the non-metric nature of the measurement space. Sethi and Chatterjee [71] use concepts of *prime events* to come up with an efficient approach to the decision tree design. Even though their method does not guarantee optimality, in most cases it provides close to optimal results, but most importantly, it is efficient.

C. Feature Selection in DTC’s

Regardless of its importance, the problem of feature subset selection in DTC’s has received relatively little attention. Kulkarni and Kanal [46] proposed the use of a branch-and-bound technique [51] to assign features to the nodes in a sequential top–down manner. The probability of error is calculated for each possible feature subset selected; if it exceeds the pre-assigned lower bound then that feature subset is abandoned.

Due to computation time, as a practical matter, the size of the feature subset to be used at each node is usually limited to be much smaller than the total number of available features [97]. Once it has been determined (often heuristically) how many features are to be used at each node, different feature subsets are examined and the one that provides maximum separability, measured by some

distance function, among classes accessible from that node, is usually selected. Note that when the number of classes is large, an exhaustive search may be unfeasible. In that case, usually a greedy-type search algorithm is used. Again, this at best would provide local optimality. Some studies, in order to determine the significance of each feature in discriminating among classes associated with each node, have offered methods that utilizes only one feature at each node. It should be noted that while the use of single feature decisions at every internal node reduces the computational burden at the tree design time, it usually leads to larger trees.

D. Decision Trees and Neural Networks

Recently, the subject of neural networks, has become very popular in many areas such as signal processing and pattern recognition. Actually, the applications of neural networks in pattern recognition problems go back to the days of simple perceptrons in the 1950's. Many advantages for neural networks have been cited in the literature (cf. Rumelhart *et al.* [67] and Lippmann [55]). The most important of them, for the pattern recognition problems, seems to be the fact that neural network-based approaches are usually nonparametric even though statistical information could be possibly incorporated to improve their performance (speed of convergence, etc.). Also neural networks can extract nonlinear combinations of features, and the resulting discriminating surfaces can be very complex. These characteristics of neural networks can be very attractive in a decision tree classifier where one has to determine the appropriate feature subsets and the decision rules at each internal node. There are various neural network models such as Hopfield nets, the Boltzmann machine, and Kohonen self-organizing feature maps, to name a few; the most popular network by far, however, is the multilayer feedforward network [67].

A multilayer feedforward neural network is simply an acyclic directed graph consisting of several layers of simple processing elements known as neurons. The neurons in every layer are fully connected to the neurons in the proceeding layer. Every neuron sums its *weighted* inputs and passes it through some kind of nonlinearity, usually a sigmoidal function. The first layer is the *input layer*, the last layer is the *output layer* and the intermediate layers are known as the *hidden layers*. See Fig. 6. Then for a supervised learning, training samples are assigned to the inputs of the network, and the connection weights between neurons are adjusted such that the error between the actual outputs of the network and the desired outputs are minimized. A commonly used error criterion is the total sum-of-squared error between the actual outputs and the desired outputs. There are many learning algorithms to perform this learning task; a frequently used learning rule is the back propagation algorithm that is basically a (noisy or stochastic) gradient descent optimization method. Note that, once one fixes the structure of the network, i.e., the number of hidden layers and the number of neurons in each hidden layer are chosen, then the network adjust its weight via the learning rule until the optimum weights are obtained. The corresponding weights (along with the structure of the network) create the decision boundaries in the feature space. The question of how to choose the structure of the network is beyond the scope of this paper and is a current research issue in neural networks. It suffices to mention that for most practical applications, networks with only one hidden layer are utilized. Also see comments in Section V and [55].

How can neural networks assist us in the design of decision trees? Gelfand and Guo [31] propose the following. Recall that in every internal node t of the tree T one would like to find the optimum (here, in the local sense) feature subset and the corresponding decision surface to partition $C(t)$, the set of classes in t , into $C(t_L)$ and $C(t_R)$, two subset of classes at nodes t_L and t_R , respectively.




Fig. 6. A three layer feedforward network with 6 inputs and 3 outputs. α_{ij} and β_{kl} are the connection weights between neuron i of input layer and neuron j of hidden layer and neuron k of hidden layer and neuron l of output layer, respectively.

That is, there are two nested optimization tasks involved here. The outer optimization loop searches for the optimum partitioning of classes into two (possibly disjoint) subsets of classes and the inner optimization loop searches for the optimum decision surface in the feature space to perform the outer optimization task. Suppose for now that we know the partitioning of $C(t)$ into $C(t_L)$ and $C(t_R)$, and are only looking for the (optimum) decision surface in the feature space. This, of course, can be easily implemented by a simple multilayer feedforward network. Then the remaining question is how to partition $C(t)$ into $C(t_L)$ and $C(t_R)$? As mentioned in Section IV-A, there are various splitting rules. One such possible rule is the impurity reduction criterion used in CART. That is, one will consider various partitioning of the $C(t)$ into $C(t_L)$ and $C(t_R)$, and choose the partitioning that gives the maximum impurity reduction (see Section IV-A and [7] for details). Of course, when the number of classes is large, an exhaustive search may be impractical. In this case, some type of greedy search can be performed.

E. Decision Rules and Search Strategies in DTC's

Once the tree structure is designed and the feature subsets to be used at each node of the tree are selected, a decision rule is needed for each node. These decision rules could be designed such that optimal performance (in any specified sense) could be attained at each node (i.e., local optimality); or the overall performance of the tree could be optimized (i.e., global optimality). Obviously, decision strategies designed to provide optimum performance at each node, do not necessarily provide overall optimum performance. For globally optimum performance, decisions made at each node should "emphasize the decision that leads to a greater joint probability of correct classification at the next level" [50]; i.e., decisions made at the different nodes are not independent. Kurzynski [49], [50] addresses both locally and globally optimum decision rules. He shows that the decision rule that provides minimum probability of error at each node (i.e., local optimality) is just the well-known maximum *a posteriori* probability rule. Since for globally optimum rules, error recognition information in the future nodes are needed, the problem is worked in a sequentially backward (bottom-up) manner starting from the terminal nodes.

Let us assume that there are m pattern classes and λ_{ij} represents the cost of losses incurred in classifying a sample from class C_i to class C_j . Then the Bayes minimum risk classifier allocates an unknown pattern sample to class C_k if

$$k = \arg \left(\min_j \sum_{i=1}^m p(C_i/X) \lambda_{ij} \right) \quad (19)$$

where X is the feature measurement vector of the unknown sample. For a "0-1" loss function, $\lambda_{ij} = 1 - \delta_{ij}$ where δ_{ij} is the Kronecker delta. Then, the previous rule becomes the simple maximum *a posteriori* rule, i.e.,

$$C_k = \arg \left(\max_i p(C_i/X) \right). \quad (20)$$

Even though "0-1" is a reasonable and frequently used loss function, in a DTC, it should only be applied at the terminal nodes where the actual classification task is performed. At the intermediate nodes in a tree, *decisions* are made as to which of the possible branches leaving that node should be followed. Let α_{ij} represent the cost incurred when a sample of class C_i is routed through branch j . Obviously, α_{ij} is a *random variable* since (Dattatreya and Kanal [13]) actions at the levels below the present level are not known. Dattatreya and Kanal [13] propose an interesting *unsupervised* scheme that *adaptively* learns the *mean* values of these random variables and improves the tree performance.

The only cost considered previously was the cost of decision making and classification. What about the cost of feature measurements? In many applications, such as medical diagnosis, feature measurement costs (e.g., laboratory tests, X-rays, etc.) may be a major portion of the cost. Thus more generally, the *total* cost should be (Dattatreya and Sarma [16], and Kulkarni and Kanal [45]) the sum of the feature measurement cost and the classification cost. Assuming that feature measurement costs are constant, Wald [80] shows that the minimum cost classifier is a multistage scheme that measures only one feature at every stage and decides the next course of action. Furthermore, Dattatreya and Sarma [16] show that rejection of class labels at intermediate stages as unlikely candidates is suboptimal in the Bayes minimum cost sense.

Now consider the situation where the tree is designed and feature subsets to be used at each node are determined. In this case, the problem of hierarchical classification can also be viewed as a problem of tree search. Kanal [40], and Kulkarni and Kanal [46] have generalized the idea of state space representation and ordered admissible search methods of Artificial Intelligence (Nilsson [60]) to the problem of hierarchical classification. Then, to search through the tree, an evaluation function $f(n)$ (Kanal [40]) is defined at each node n as

$$f(n) = g(n) + h(n) + l(n) \quad (21)$$

where $g(n)$ is the cost from the root to node n , $h(n)$ is (an estimate of) the cost from n to a terminal node *accessible* from n , and $l(n)$ is (an estimate of) the risk of classification at a terminal node accessible from n .

At a goal node, s^* , the total cost is the sum of measurement costs at each node along the path from the root to s^* plus the actual risk, $r(s^*)$, associated with s^* . If $r(s^*)$ is estimated based on only the feature measurements at each node along the path from the root to s^* , the strategy that provides minimum total cost is called an *S-admissible* search strategy. When $r(s^*)$ is estimated using all the measurements, not just those on the path to s^* , then the strategy that provides minimum total cost is known as a *B-admissible* search strategy (Kanal [40]).

At every node, a lower bound on the value of risk is used in the evaluation function for the *S-admissible* search, except for a goal node where the actual risk is used. In a *B-admissible* search, however, since the risk associated with any goal s^* can change as additional measurements are made along other paths in the tree, a search could not be terminated simply as one goal state is reached. Therefore, an upper bound function on the risk is also needed to determine the termination point (Kanal [40], Kulkarni and Kanal

[46]). Obviously, the performance of either of the two aforementioned search methods is dependent on one's ability to find proper bounds on the goal risk. The advantage of the state-space model as pointed out by Kulkarni and Kanal [46] is that, in contrast to the hierarchical classifier approach where only one path is pursued from the root to a terminal point, here provisions are made to back up and follow other alternative paths in the tree if it is needed. Obviously, however, this would increase the number of nodes visited before a decision is made and thus reduce the efficiency. Also search efficiency is a function of the tightness of the bounds on the risk; and tight bounds are difficult to obtain particularly for some continuous distributions [46]. It is also interesting to note that the way cost functions are usually defined neither the *B-admissible* nor the *S-admissible* search has provisions for the time component [81].

As mentioned earlier, when the number of classes is large, i.e., on the order of hundreds or even thousands as is in the Chinese character recognition problem for instance, tree classifiers could provide a considerable amount of time savings. But error has a tendency to accumulate from level to level. This is because in problems with large numbers of classes, the tree has many levels. If e is the minimum error at each node of each level and there are on the average l levels in the tree, the average total error rate of the tree would be on the order of $(l * e)$. As proposed in Wang and Suen [81], there are two ways to reduce the total error: either reduce e or have a better search strategy capable of back-up and rerouting. In general e can only be reduced by allowing overlaps at internal nodes [95]; but this would increase the number of nodes and terminal nodes, thus greatly increasing the memory space requirements. To improve tree search, however, several heuristic search strategies are proposed [2], [9], [10], [28], [33], [35], [81]–[83]. Chang and Pavlidis [10] proposed a branch-bound-backtracking algorithm. A fuzzy decision function [99], taking values in the interval $[0,1]$ is used to assign the decision values of all the branches going out of a node. The overall decision value of a path is defined to be the product of the decision values of all the branches along that path. The product operator could be either the regular product or the max-min operator [9]. Branches with the largest decision values are followed. At any node, if the total decision value up to that node falls below a prespecified threshold, the algorithm backtracks to the previous node(s) and follows another path. Of course, the speed of the search is related to the threshold value set.

Since the most confusion and error occurs for samples that come from regions near the boundaries between classes (or groups of classes), Wang and Suen [81] propose a two-step search approach. In the first step, following the usual top-down search approach, an unknown sample reaches a terminal node. The degree of similarity of the sample with the class label associated with that node is computed. If this value is above a preset threshold, the sample is classified and the next sample is treated; otherwise, the unknown sample goes through the second search called "fuzzy logic search." As argued in [81], the rationale for fuzzy logic is that somehow the global search information and all the possible terminals (classes) must be recorded. Probabilistic decision values based on Bayes decision regions, although perhaps more precise, are usually much harder to estimate. Fuzzy decision values, however, are both flexible and easier to compute. By adjusting the two thresholds involved in two steps, the speed of the search could be controlled. This would also provide a tradeoff between speed and accuracy of the tree classifier.

V. OTHER TREE RELATED ISSUES

A. Incremental Tree Design

With regard to the training samples, there are two possibilities:

- 1) The *entire* training sample is available at the time of decision tree

design; 2) the training sample arrives in a *stream*.

For case 1, once a tree is designed the task is completed and such design algorithms are known as non-incremental algorithms. For the second case, however, one has two options:

- a) whenever the new training samples become available, discard the current tree and construct a replacement tree using the enlarged training set, or
- b) revise the existing tree based on the new available information.

Procedures corresponding to b) are known as *incremental* decision tree design algorithms. Of course, it would be desirable that the incremental decision tree design algorithms produce the *same* trees as those if all the training samples were available at the time of design. Utgoff [79] has developed one such incremental algorithm.

B. Tree Generalization

In designing a DTC, one must always keep in mind that the tree designed is going to be used to classify unseen test samples. With this in mind, one has two options: 1) design a decision tree that correctly classifies *all* the training samples, also known as a *perfect tree* [62], [63]; and select the *smallest* perfect tree; or 2) construct a tree that is perhaps imperfect (in the previous sense) but has the smallest possible error rate in classification of test samples. In practical pattern recognition tasks, it is usually this second type of tree that is of most interest.

Regardless of which of the previous two types of trees one may need, it is usually desirable to keep the size of the tree as small as possible. The reasons for this are: 1) smaller trees are more efficient both in terms of tree storage requirements and test time requirements and 2) smaller trees tend to generalize better for the unseen samples because they are less sensitive to the statistical irregularities and idiosyncrasies of the training data.

C. Missing Value Problem

The missing value problem can occur in either the design phase, or the test phase, or both. In the design step, suppose some of the training sample feature vectors are incomplete; that is some feature elements of some feature vectors are not recorded or are missing. This can happen, for instance, due to some occasional sensor failures. Similarly for the test samples, some feature values may be missing.

In the design phase, one simple but wasteful method to cope with this problem is to throw away the incomplete feature vectors. For the test sample, of course, this simple option is not acceptable. Breiman *et al.* [7] propose the following solution that is based on the idea of using *surrogate splits*. For the case of simplicity, consider the case of binary splitting at every internal node based on the value of only *one* feature. For the extension of this idea to a linear combination of features, see [7] and [21]. The basic idea is as follows.

In the tree design phase, at node t , find the best split S_m^* on the feature element X_m using all the training samples containing a value of X_m . Then select the split S^* , which maximizes the impurity reduction $\Delta i(S_m^*, t)$ at node t (see Section IV-A).

For the incomplete test sample, if at a node t the best split S^* is not defined because of missing feature element values, proceed as follows. Define a *measure of similarity* between two splits S_i and S_j as $\lambda(S_i; S_j)$. Examine all non-missing feature elements for the test sample; find that one, say X_m , with split \hat{S}_m , that is most similar to S^* . \hat{S}_m is called a *surrogate split* to S^* . Then use \hat{S}_m at node t to decide to traverse to node t_L or t_R .

D. Robustness of Decision Tree Design

Since decision trees are often constructed by using just some sets of training samples, it is important to make sure that the design

procedure is in some sense robust relative to the presence of "bad" samples or outliers. Of course, one could always edit the training data before the application in the tree design. This, however, in many cases may not be feasible for the following reasons.

- 1) The designer may not be aware of the existence of outliers at the time of the tree design, even though outliers are usually easily detected in a training set; in some cases a thorough examination of data set may be necessary.
- 2) With a small sample, one may not want to throw away any valuable samples.

E. Decision Trees and Neural Networks

Even though multilayer perceptron (MLP) networks and DTC's are two very different techniques for classification, the similarities in the distributed nature of decision-making in both processes has motivated some researchers [5] to perform some empirical comparative studies. Following are some of the general conclusions drawn so far.

- 1) Both DTC's and MLP networks are capable of generating arbitrarily complex decision boundaries from a given set of training samples; usually neither needs to make any assumptions about the distributions of the underlying processes.
- 2) MLP networks are usually more easily capable of providing incremental learning than DTC's.
- 3) Training time for MLP networks is, however, usually much longer than the training time for DTC's. Therefore, MLP networks tend to be most useful in the circumstances where the training is infrequent and the given classifier is used for a long period.
- 4) In the empirical studies reported in [5], MLP networks performed as well as or better than their CART counterpart.
- 5) DTC's are sequential in nature, compared to the massive parallelism usually present in MLP networks.
- 6) The behavior of DTC's in general, and CART in particular, make it much more useful for data interpretation than MLP networks. An MLP network, in general, is no more than a black box of weights; and reveals little information about the problem it's applied to.
- 7) Perhaps the most interesting conclusion in [5] is that there is not enough evidence (theoretical or empirical) to provide a strong support for either one of the approaches alone.

Then, perhaps a more useful stand to take with regards to the decision tree and neural networks is to consider both of these tools together. This may be looked at from following two perspectives.

- 1) Neural networks may be used in the implementation of various aspects of decision trees. For instance, Gelfand and Guo [31] use neural networks in the internal nodes of a decision tree to perform the task of feature selection and decision boundary construction. In this approach, one is constructing a *tree of neural networks*. See also Golea and Marchand [27].
- 2) The alternate possibility is to design the DTC and then map the tree into a neural network. This is the direction Sethi [69]–[70], and Sethi and Otten [72] have taken. The main idea here is that, as mentioned earlier in Section IV-C, in the design of multilayer feedforward networks, the structure of the network, i.e., the number of hidden layers and the number of neurons in each hidden layer is not known in advance, and is often chosen rather heuristically and by trial and error. The method described in [69], [70], and [72] offers one way to uncover the structure of the network. That is, they offer a method to implement a predesigned decision tree via a multilayer feedforward neural network. Perhaps one of the main advantages of this approach, other than possibly eliminating guess work in the design of

MLP networks, is that the network designed this way usually has fewer connections. This may be an important factor from VLSI implementation point of view.

VI. CONCLUSION

DTC's show a great deal of potential in many pattern recognition problems such as remotely sensed multisource data classification, medical diagnosis, speech, and character recognition, to mention a few. Perhaps one of the main features of DTC's is the flexibility they provide; for example, the capability of using different feature subsets and decision rules at different stages of classification and the capability of tradeoffs between classification accuracy and time/space efficiency.

Some of the goals of this review have been to:

- 1) bring the disparate issues in decision tree classifiers closer together and perhaps motivate some new ideas;
- 2) provide a more unified view of decision tree classifiers;
- 3) and caution the "casual" users of these methods of the possible "pitfalls" of each method.

Even though optimization techniques such as dynamic programming have been offered by some researchers to define trees that have certain optimal characteristics, a *practical* and *feasible* solution that is truly optimal with respect to the choice of tree structure, feature subsets, and decision rule strategies is yet far from realization. This is mainly due to either unavailability of the necessary information (e.g., *a priori* and class conditional statistics) for the design or to complexities of the proposed methods that limits their usefulness for real applications of moderate or large size. Recognizing the difficulties involved in the

design of optimal DTC's, many heuristic methods have been proposed in the literature.

The four basic approaches for design of DTC's examined in this review are as follows.

- 1) Bottom-up approach where one starts with the information classes and continues combining classes until one is left with a node containing all the classes (i.e., the root).
- 2) Top-down approach where starting from the root node, using a splitting rule, classes are divided until a stopping criterion is met. The main issues in this approach are: a) choice of splitting criterion; b) stopping rules; c) labeling the terminal nodes.
- 3) Hybrid approach where one uses a bottom-up procedure to direct and assist a top-down procedure.
- 4) Tree growing-pruning approach where in order to avoid some difficulties in choosing a stopping rule, one grows the tree to its maximum size where the terminal nodes are pure or almost pure, and then selectively prunes the tree.

As was pointed out earlier, computational efficiency, accuracy of classification and storage space requirements usually seem to be conflicting requirements. This may actually assist the designer (user) of a DTC in selecting one possible scheme over the others. The existing methods with their relative time/space complexities, degrees of (sub)optimality and proposed areas of applications are summarized in the following tables. Table I provides a summary of some of the tree design methods in terms of the assumptions each approach makes, their performance criterion and some of the specific requirements of each method. Table II gives the brief summary of various feature selection methods. Table III summarizes the various decision rule and search strategies used in decision trees. And finally, Table IV provides specific references for various applications.

TABLE I
SUMMARY OF TREE STRUCTURE DESIGN

Ref.	Optimal?	Criterion	Assumptions	Remarks (by author or alternate reference)
[4]	locally (node-wise)	prob. of error	<ul style="list-style-type: none"> given: a well defined feature transformation 	<ul style="list-style-type: none"> an automated technique; uses only <i>a priori</i> statistics; a procedure to compute the global P_e is also provided.
[8]	no	Max. inf. gain at each node	<ul style="list-style-type: none"> given: a binary image with <i>a priori</i> class prob.'s; number of black bits in each position in each class in the design samples. 	<ul style="list-style-type: none"> input (for classification) is a 2-D binary pattern computationally efficient; storage requirements reasonable; especially suited for (optical) character recognition.
[15]	yes	Min. total cost (feature meas. cost + classif. cost)	<ul style="list-style-type: none"> Design samples available; conditional densities are unimodal. 	<ul style="list-style-type: none"> provides a binary tree; superior performance to the method proposed in ref. #13; both in terms of simplicity of design and the total expected cost of classification
[16]	no	"	—	<ul style="list-style-type: none"> heuristic; provides a binary tree.

Continued on next page

TABLE I (Continued)
SUMMARY OF TREE STRUCTURE DESIGN

Ref.	Optimal?	Criterion	Assumptions	Remarks (by author or alternate reference)
45	yes	Weighted sum of P_e and avg. meas. cost in classifying a random sample	<ul style="list-style-type: none"> • Uses ML at each node • overall P_e of tree is small; • features on any path from root to a terminal node are statist. indep. 	<ul style="list-style-type: none"> • a bottom-up approach; • time and space complexities are high.
52	no	—	<ul style="list-style-type: none"> • training sample avail. 	<ul style="list-style-type: none"> • unsupervised clustering of <i>a priori</i> classes
53	no	accuracy	<ul style="list-style-type: none"> • labeled training pattern available. 	<ul style="list-style-type: none"> • examines one feature at each node; • finds best feature and corresp. threshold for each node.
57	yes	a general cost function.	<ul style="list-style-type: none"> • given: samples from all the classes. 	<ul style="list-style-type: none"> • binary tree; • storage requirements may be large.
61	yes	"	"	<ul style="list-style-type: none"> • single feature comparison at each node; • feasible for small # of features; • given a specific partitioning of the feature space (i.e., P_e fixed), this method could be used to obtain a tree with, for instance, min. # of nodes.
65	locally	min. Bayes error.	<ul style="list-style-type: none"> • two-class problem; • training samples available. 	<ul style="list-style-type: none"> • integrates binary tree construction and feature selection problem; • efficient for 2-class problem; • can evaluate contribution of each feature to the classification task.
68	yes	prob. of error	<ul style="list-style-type: none"> • all the N features are avail. at every node 	<ul style="list-style-type: none"> • estimates a posteriori prob.'s, but in a sequence of steps; • asymptotically robust w.r.t. changes in tree structure.
71	no	<ul style="list-style-type: none"> • min. expected cost of a tree; • weights of the terminal nodes and meas. cost 	<ul style="list-style-type: none"> • meas. costs are known; • each meas. can be used only once in decision making. 	<ul style="list-style-type: none"> • useful for discrete feature applications; • efficient; • easy to implement.
73	locally	max. mutual inf. gain at each node.	<ul style="list-style-type: none"> • given: a set of design samples and a pre-spec. desired P_e. 	<ul style="list-style-type: none"> • examines 1 feature per node; • provides a binary structure; • easy to implement; • feature selection capability inherent in the algorithm.
76	yes	max. entropy reduction and min. overlap.	—	<ul style="list-style-type: none"> • a clustering algorithm; • useful for large # of classes
77	sub-opt.	weighted sum of accuracy & efficiency.	<ul style="list-style-type: none"> • decision rule at each stage is Gaussian ML 	<ul style="list-style-type: none"> • many applications in remote sensing explained here.
95	yes	accuracy, or total perform	<ul style="list-style-type: none"> • uses training set to estimate appropriate densities 	—
97	no	efficiency and accuracy	<ul style="list-style-type: none"> • multivariate Gaussian dist. for all the classes 	<ul style="list-style-type: none"> • uses linear discriminant function to improve speed; • number of features at each node is pre-specified.

TABLE II
SUMMARY OF FEATURE SELECTION

Reference	Optimal	Criterion	Remarks (by author or alternate reference)
[45]	yes	accuracy	Not very efficient
[53]	no	accuracy	Finds one best feature at each node
[65]	locally	accuracy	"
[68]	yes	max. accuracy in estimating proper <i>a posteriori</i> prob.'s	—
73]	locally	max. inf. gain at each node.	Finds one best feature at each node.
[52], [77], [95] [97]	no	max. separability between classes at each node, usually measured by some distance function.	Usually by some heuristic methods number of features to be used at each node are specified; most practical (feasible) ones.

TABLE III
SUMMARY OF DECISION RULE AND SEARCH STRATEGIES

Reference	Optimal?	Criterion?	Remarks
[40]	no	min. total cost of decision making; i.e., sum of costs at different nodes plus risk of final classification.	• several search methods within the context of pattern classification are addressed here.
[45]	yes	weighted sum or correct recognition rate and avg. meas. cost incurred per sample.	• assumes knowledge of all joint class conditional probabilities; • some methods to reduce computational complexities are offered.
[46]	no	—	• crucial task is to find tight bounds on the goal risk; • able to back-up & re-route, thus can improve correct recogn. rate.
[49], [50]	yes	probability of error	• assumes: tree structure and feature subsets to be used at each node are given.
[81]	no	allow trade-off between accuracy & efficiency.	• uses: fuzzy logic search to improve speed, and global training to improve correct recogn. rate.

TABLE IV
CITED APPLICATIONS

Reference	Cited Application
[4], [6], [41], [77], [95], [97]	remote sensing
[52], [58], [98]	medical
[15], [16]	speech
[8], [9], [10], [30], [53], [68], [71], [73], [81]–[84]	character recognition
[7], [40], [45]–[48] [57], [61]	general

REFERENCES

- [1] J. Aczel and J. Daroczy, *On Measures of Information and Their Characterizations*. New York: Academic, 1975.
- [2] R. Ahlsmede and I. Wegeru, *Search Problems*. New York: Wiley-Interscience, 1987.
- [3] A. V. Aho, J. E. Hopcroft, and J. D. Ullmann, "The Design and Analysis of Computer Algorithms". Reading, MA: Addison-Wesley, 1974.
- [4] P. Argentiero, R. Chin, and P. Beaudet, "An automated approach to the design of decision tree classifiers," *IEEE Trans. Patt. Anal. Mach. Intell.*, vol. PAMI-4, pp. 51–57, 1982.
- [5] L. Atlas *et al.*, "A performance comparison of trained multilayer perceptrons and trained classification trees," *Proc. IEEE*, vol. 78, pp. 1614–1619, 1990.
- [6] L. A. Bartolucci, P. H. Swain, and C. Wu, "Selective radiant temperature mapping using a layered classifier," *IEEE Trans. Geosci. Electron.*, vol. GE-14, pp. 101–106, 1976.
- [7] L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone, *Classification and Regression Trees*. Belmont, CA: Wadsworth Int., 1984.
- [8] R. C. Casey and G. Nagy, "Decision tree design using a probabilistic model," *IEEE Trans. Inform. Theory*, vol. IT-30, pp. 93–99, 1984.
- [9] R. L. P. Chang, "Application of fuzzy decision techniques to pattern recognition and curve fitting," Ph.D. dissertation, Dep. of EECS, Princeton Univ., Princeton, NJ, 1976.

Finally, it should be pointed out that since the performance of a DTC depends on how well the tree is designed, special attention should be paid to this phase. In particular, when the number of test samples to be classified is large, as is, for example, in remote sensing applications, the time spent for the design could be well justified.

- [10] R. L. P. Chang and T. Pavlidis, "Fuzzy decision tree algorithms," *IEEE Trans. Syst., Man, Cybern.*, vol. SMC-7, pp. 28–35, 1977.
- [11] P. A. Chou and R. M. Gray, "On decision trees for pattern recognition," in *Proc. IEEE Symp. Inform. Theory*, Ann Arbor, MI, p. 69, 1986.
- [12] M. Chou, T. Lookabaugh, and R. M. Gray, "Optimal pruning with applications to tree structured source coding and modeling," *IEEE Trans. Inform. Theory*, vol. T35, pp. 299–315, 1989.
- [13] G. R. Dattatreya and L. N. Kanal, "Adaptive pattern recognition with random costs and its application to decision trees," *IEEE Trans. Syst., Man, Cybern.*, vol. SMC-16, pp. 208–218, 1986.
- [14] G. R. Dattatreya and L. N. Kanal, "Decision trees in pattern recognition," in *Progress in Pattern Recognition 2*, Kanal and Rosenfeld, Eds., Elsevier Science, pp. 189–239, 1985.
- [15] G. R. Dattatreya and V. V. S. Sarma, "Decision tree design for pattern recognition including feature measurement cost," in *Proc. 5th Int. Conf. Pattern Recognition*, vol. II, pp. 1212–1214, 1980.
- [16] G. R. Dattatreya and V. V. S. Sarma, "Bayesian and decision tree approaches for pattern recognition including feature measurement costs," *IEEE Trans. Patt. Anal. Mach. Intell.*, vol. PAMI-3, pp. 293–298, 1981.
- [17] E. Diday and J. V. Moreau, "Learning hierarchical clustering from examples—Applications to the adaptive construction of dissimilarity indices," *Pattern Recognition Lett.*, pp. 223–230, 1986.
- [18] R. O. Duda and P. E. Hart, *Pattern Classification and Scene Analysis*. New York: Wiley-Interscience, 1973.
- [19] P. Fletcher and M. J. D. Powell, "A rapid decent method for minimization," *Computer J.*, vol. 6, ISS.2, pp. 163–168, 1963.
- [20] J. H. Friedman, "A variable metric decision rule for nonparametric classification," Stanford Linear Accelerator Center-PUB-1573, CS-75-487, Apr. 1975.
- [21] —, "A recursive partitioning decision rule for nonparametric classifier," *IEEE Trans. Comput.*, vol. C-26, pp. 404–408, 1977.
- [22] K. S. Fu, *Sequential Methods in Pattern Recognition and Machine Learning*. New York: Academic, 1968.
- [23] K. Fukunaga, *Introduction to Statistical Pattern Recognition*. New York: Academic, 1972.
- [24] K. Fukunaga and P. M. Narendra, "A branch and bound algorithm for computing k -nearest neighbors," *IEEE Trans. Comput.*, vol. C-24, pp. 750–753, 1975.
- [25] S. B. Gelfand, C. S. Ravishanker, and E. J. Delp, "An iterative growing and pruning algorithm for classification tree design," *IEEE Trans. Patt. Anal. Mach. Intell.*, pp. 163–174, 1991.
- [26] B. V. Gnedenko and V. S. Koreljuk, "On the maximum discrepancy between two empirical distributions," *Dokl. Acad. Nauk.*, vol. SSSR 4, pp. 525–528, 1951.
- [27] M. Golea and M. Marchand, "A growth algorithm for neural networks," *Europhys. Lett.*, vol. 12, no. 3, pp. 205–210, 1990.
- [28] S. W. Golomb and L. D. Bavmert, "Backtrack programming," *J. Ass. Comput. Mach.*, vol. 12, pp. 516–524, 1965.
- [29] R. M. Goodman and P. Smyth, "Decision tree design from a communication theory standpoint," *IEEE Trans. Inform. Theory*, vol. IT-34, pp. 979–994, 1988.
- [30] Y. X. Gu, Q. R. Wang, and C. Y. Suen, "Application of multi-layer decision tree in a computer recognition of Chinese characters," *IEEE Trans. Patt. Anal. Mach. Intell.*, vol. PAMI-5, pp. 83–89, 1983.
- [31] S. Gelfand and H. Guo, "Tree classifiers with multilayer perceptron feature extraction," Ph. D. dissertation, School of Elec. Eng., Purdue Univ., W. Lafayette, IN, May 1991.
- [32] D. E. Gustafson, S. B. Gelfand, and S. K. Mitter, "A nonparametric multiclass partitioning methods for classification," in *Proc. 5th Int. conf. pattern Recognition*, pp. 654–659, 1980.
- [33] P. A. V. Hall, "Branch-and-Bound and Beyond," in *Proc. 2nd Joint Int. Conf. Artificial Intelligence*, 1971.
- [34] R. M. Haralick, "The table look-up rule," in *Proc. Conf. Pattern Recognition*, p. 447, 1976.
- [35] P. E. Hart, "Searching probabilistic decision trees," AI Group Tech. Note No. 2, SRI project 7494, SIR, Stanford, CA, 1969.
- [36] C. R. P. Hartmann, P. K. Varshney, K. G. Mehrotra, and C. L. Gerberich, "Application of information theory to the construction of efficient decision trees," *IEEE Trans. Inform. Theory*, vol. IT-28, pp. 565–577, 1982.
- [37] E. G. Henrichon, Jr. and K. S. Fu, "A nonparametric partitioning procedure for pattern classification," *IEEE Trans. Comput.*, vol. C-18, pp. 604–624, 1969.
- [38] G. F. Hughes, "On the mean accuracy of statistical pattern recognizers," *IEEE Trans. Inform. Theory*, vol. IT-14, pp. 55–63, 1968.
- [39] L. Hyafil and R. L. Rivest, "Constructing optimal binary decision trees is NP-complete," *Inform. Processing Lett.*, vol. 5, no. 1, pp. 15–17, 1976.
- [40] L. N. Kanal, "Problem-solving methods and search strategies for pattern recognition," *IEEE Trans. Patt. Anal. Mach. Intell.*, vol. PAMI-1, pp. 193–201, 1979.
- [41] B. Kim and D. A. Landgrebe, "Hierarchical decision tree classifiers in high-dimensional and large class data," Ph.D. dissertation and Tech. Rep. TR-EE-90-47, School of Elec. Eng., Purdue Univ., W. Lafayette, IN, 1990.
- [42] P. R. Krishnaiah, Ed., *On Hierarchical Classifier and Interactive Design, in Applications of Statistics*. Amsterdam, The Netherlands: North Holland, 1971, pp. 301–321.
- [43] D. E. Knuth, "Optimum binary search trees," *ACTA Inform.*, vol. 1, pp. 14–25, 1971.
- [44] D. E. Knuth, "The art of computer programming, 1: Fundamental algorithms," Reading, MA: Addison-Wesley, 1968.
- [45] A. V. Kulkarni and L. N. Kanal, "An optimization approach to hierarchical classifier design," in *Proc. 3rd Int. Joint Conf. Pattern Recognition*, San Diego, CA, 1976.
- [46] A. V. Kulkarni and L. N. Kanal, "Admissible search strategies for parametric and non-parametric hierarchical classifiers," in *Proc. 4th Int. Conf. Pattern Recognition*, Kyoto, Japan, 1978.
- [47] A. V. Kulkarni, "On the mean accuracy of hierarchical classifiers," *IEEE Trans. Comput.*, vol. C-27, pp. 771–776, 1978.
- [48] A. V. Kulkarni, "Optimal and heuristic synthesis of hierarchical classifiers," Ph.D. dissertation, Univ. Maryland, College Park, Comput. Sci. Tech. Rep. TR-469, 1976.
- [49] M. W. Kurzynski, "Decision rules for a hierarchical classifier," *Pattern Recog. Lett.*, vol. 1, pp. 305–310, 1983.
- [50] M. W. Kurzynski, "The optimal strategy of a tree classifier," *Pattern Recog.*, vol. 16, pp. 81–87, 1983.
- [51] E. L. Lawler and D. E. Wood, "Branch-and-bound: A survey," *Operation Res.*, vol. 14, 1966.
- [52] G. Landeweerd, T. Timmers, E. Gersema, M. Bins and M. Halic, "Binary tree versus single level tree classification of white blood cells," *Pattern Recog.*, vol. 16, pp. 571–577, 1983.
- [53] X. Li and R. C. Dubes, "Tree classifier design with a Permutation statistic," *Pattern Recog.*, vol. 19, pp. 229–235, 1986.
- [54] Y. K. Lin and K. S. Fu, "Automatic classification of cervical cell using a binary tree classifier," *Pattern Recog.*, vol. 16, pp. 69–80, 1983.
- [55] R. P. Lippmann, "An introduction to computing with neural nets," *IEEE ASSP Mag.*, pp. 4–22, Apr. 1987.
- [56] C. McMillan, Jr., *Mathematical Programming*. New York: Wiley, 1975, ch. 8.
- [57] W. S. Meisel and D. A. Michalopoulos, "A Partitioning algorithm with application in pattern classification and optimization of decision trees," *IEEE Trans. Comput.*, vol. C-22, pp. 93–103, Jan. 1973.
- [58] J. Mui and K. S. Fu, "Automated classification of nucleated blood cells using a binary tree classifier," *IEEE Trans. Patt. Anal. Mach. Intell.*, vol. PAMI-2, pp. 429–443, 1980.
- [59] M. Nadler, "Error and reject rates in a hierarchical pattern recognizer," *IEEE Trans. Comput.*, vol. C-19, pp. 1598–1601, 1970.
- [60] N. J. Nilsson, *Problem-Solving Methods in Artificial Intelligence*. New York: McGraw-Hill, 1971.
- [61] H. Payne and W. Meisel, "An algorithm for constructing optimal binary decision trees," *IEEE Trans. Comput.*, vol. C-26, pp. 905–916, 1977.
- [62] J. R. Quinlan, "Induction of decision trees," *Machine Learning*, vol. 1, pp. 81–106, 1986.
- [63] —, "Decision trees and decision-making," *IEEE Trans. Syst., Man, Cybern.*, vol. SMC-20, pp. 339–346, 1990.
- [64] J. R. Quinlan and R. L. Rivest, "Inferring decision trees using the minimum description length principle," *Information and Computation*, vol. 80, no. 3, pp. 227–248, 1989.
- [65] E. Rounds, "A combined non-parametric approach to feature selection and binary decision tree design," *Pattern Recog.*, vol. 12, pp. 313–317, 1980.
- [66] E. Rounds, "Computation of two-sample Kolmogorov–Smirnov statistic," Project Memo TSC-PM-A 142-5, Technology Service Corp., Santa Monica, CA, 1978.
- [67] D. E. Rumelhart and J. L. McClelland, Eds., *Parallel Distributed Processing*, vol. 1. Cambridge, MA: MIT Press, 1986.
- [68] J. Schuermann and W. Doster, "A decision-theoretic approach in hierarchical classifier design," *Pattern Recog.*, vol. 17, pp. 359–369, 1984.
- [69] I. K. Sethi, "Entropy nets: From decision trees to neural networks," *Proc. IEEE*, vol. 78, pp. 1605–1613, Oct. 1990.
- [70] —, "Layered neural net design through decision trees," in *Proc. Int. Symp. Circuits and Syst.*, New Orleans, LA, May 1–3, 1990.
- [71] I. K. Sethi and B. Chatterjee, "Efficient decision tree design for dis-

- crete variable pattern recognition problems," *Pattern Recog.*, vol. 19, pp. 197–206, 1977.
- [72] I. K. Sethi and M. Otten, "Comparison between entropy net and decision tree classifiers," in *Proc. Int. Joint Conf. Neural Net. IJCNN*, San Diego, CA 1989.
- [73] I. K. Sethi and G. Sarvarayudu, "Hierarchical classifier design using mutual information," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. PAMI-4, pp. 441–445, 1982.
- [74] Q. Y. Shi, "A method for the design of binary tree classifiers," in *Proc. IEEE Conf. Image Processing and Pattern Recog.*, pp. 21–26, 1981.
- [75] C. R. Smith and W. T. Grandy, Jr., eds., *Maximum-Entropy and Bayesian Methods in Inverse Problems*. The Netherlands: Reidel, 1985.
- [76] P. Y. Suen and Q. R. Wang, "ISOETRP—An interactive clustering algorithm with new objectives," *Pattern Recog.*, vol. 17, pp. 211–219, 1984.
- [77] P. Swain and H. Hauska, "The decision tree classifier design and potential," *IEEE Trans. Geosci. Electron.*, vol. GE-15, pp. 142–147, 1977.
- [78] J. L. Talmon, "A multiclass nonparametric partitioning algorithm," in E. S. Gelsema and L. N. Kanal, Eds., *Pattern Recog. in Practice II*. Amsterdam, The Netherlands: Elsevier Science, 1986.
- [79] P. E. Utgoff, "Incremental induction of decision trees," *Machine Learning*, vol. 4, pp. 161–186, 1989.
- [80] A. Wald, *Sequential Analysis*. New York: Wiley, 1947.
- [81] Q. R. Wang and C. Y. Suen, "Large tree classifier with heuristic search and global training," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. PAMI-9, pp. 91–102, 1987.
- [82] Q. R. Wang and C. Y. Suen, "Analysis and design of a decision tree based on entropy reduction and its application to Large character set recognition," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. PAMI-6, pp. 406–417, 1984.
- [83] Q. R. Wang and C. Y. Suen, "Classification of Chinese characters by phase feature and fuzzy logic search," in *Proc. Int. Conf. Chinese Inform. Processing*, vol. 1, Beijing, China, pp. 133–155, 1983.
- [84] Q. R. Wang, Y. X. Gu, and C. Y. Suen, "A preliminary study on computer recognition of Chinese characters printed in different fonts," in *Proc. Int. Conf. Chinese Lang. Comput. Soc.*, pp. 344–351, 1982.
- [85] Watanabe, "Pattern recognition as conceptual morphogenesis," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. PAMI-2, pp. 161–165, 1980.
- [86] S. Watanabe, *Knowing and Guessing*. New York: Wiley, 1969.
- [87] ———, *Pattern Recognition: Human and Machine*. New York: Wiley-Interscience, 1985.
- [88] ———, "Learning process and inverse H-theorem," *IRE Trans. Inform. Theory*, vol. IT-8, p. 246, 1962.
- [89] ———, "Information-theoretical aspects of inductive and deductive inference," *IBM J. Res. Develop.*, vol. 4, p. 208, 1960.
- [90] ———, "Application of dynamical coalescence model of clustering," in *Proc. 3rd Int. Joint Conf. Pattern Recog.*, San Diego, CA, p. 176, 1976.
- [91] ———, "Karhunen-Loeve expansion and factor analysis," in *Trans. 4th Prague Conf. Inform. Theory*, 1965.
- [92] S. Watanabe and E. T. Harada, "A dynamical model of clustering," in *Proc. 2nd Int. Joint Conf. Pattern Recog.*, Copenhagen, Denmark, p. 413, 1974.
- [93] S. Watanabe, "Pattern recognition as a quest for minimum entropy," *Pattern Recog.*, vol. 13, p. 381–387, 1981.
- [94] R. Winston, "A heuristic program that constructs decision trees," MIT project MAC, Memo #173, 1969.
- [95] C. Wu, D. Landgrebe, and P. Swain, "The decision tree approach to classification," School Elec. Eng., Purdue Univ., W. Lafayette, IN, Rep. RE-EE 75-17, 1975.
- [96] S. S. Yau and J. M. Garnet, "Least mean square approach to pattern classification," in *Frontier of Pattern Recognition*, New York: Academic, 1972, pp. 575–588.
- [97] K. C. You and K. S. Fu, "An approach to the design of a linear binary tree classifier," in *Proc. 3rd Symp. Machine Processing of Remotely Sensed Data*, Purdue Univ., W. Lafayette, IN, 1976.
- [98] S. Q. Yun and K. S. Fu, "A method for the design of binary tree classifiers," *Pattern Recog.*, vol. 16, pp. 593–603, 1983.
- [99] L. A. Zadeh, "Fuzzy sets," *Inform. Contr.*, pp. 338–353, 1965.

Character Recognition in a Sparse Distributed Memory

Young-Sik Hong and Su-Shing Chen

Abstract—Kanerva's sparse distributed memory model is applied to character recognition. The results of recognizing corrupted characters, using a sparse distributed memory with fewer than 10 000 locations, is described; 100 corrupted test patterns were generated to recognize 18 template patterns and performance evaluation of the model is obtained. Insights about the behavior of sparse distributed memories as well as the model's applicability to character recognition, are provided.

I. INTRODUCTION

The sparse distributed memory of Kanerva [4] is a mathematical model whose architecture models the human memory. In a sparse distributed memory, each pattern is represented as an n bit pattern and multiple copies of the pattern are stored over the entire memory, while a single copy of the pattern is stored in the conventional memory.

This form of data storage enables retrieval of the most similar template patterns given a corrupted input pattern. While many accurate character recognition algorithms exist already, the distributed memory model is noteworthy, because template patterns can be easily trained by a simple storing operation and moderately noised input characters can be recognized (See [3] and [5] for a somewhat different model). These properties suggest that the distributed memory model might be useful in solving complex cognition problems [1], [2], [7].

In [4], Kanerva established a memory model based on n dimensional vector space, and suggested its size in the range of 2^{100} to 2^{1000} . To overcome the requirement of a very large address space, he assigned only 1 000 000 storage locations. In our work, we have developed a simulation model of the sparse distributed memory in $\{0, 1\}^{1000}$ with fewer than 10 000 storage locations of dimension 30. This model was used to recognize several characters, with considerable noise.

II. THE SPARSE DISTRIBUTED MEMORY MODEL

In a sparse distributed memory, many storage locations participate in a single write or read operation, in contrast to the conventional computer memory, in which only one location is active at a single write or read operation. A data item is stored by distributing copies of it in many locations and retrieved on the basis of the similarity of addresses. An address decoder would select a group of locations based on the input address. For this distributed operation, linear threshold neurons are used for address decoding [4]. Each storage location has one address decoder neuron, n address lines, n output lines, and status flags, as shown in Fig. 1.

A. Address Decoding of Sparse Memory

For an arbitrary storage location L , let an n -tuple $w = \langle w_0, w_1, \dots, w_{n-1} \rangle$ be input coefficients for the address decoder neuron, and let a binary vector $x = \langle x_0, x_1, \dots, x_{n-1} \rangle$ be the address of L , where $w_i = 1$ or -1 , $0 \leq i \leq n-1$, and $x_i = 0$

Manuscript received June 18, 1990; revised December 7, 1990. This work was supported in part by the Korea Science and Engineering Foundation and in part by the National Science Foundation.

Y. S. Hong is with the Department of Computer Science, Dongguk University, Seoul 100-715, Republic of Korea.

S. S. Chen is with the Department of Computer Science, University of North Carolina, Charlotte, NC 28223.

IEEE Log Number 9042578.